

hw1__code__output

March 10, 2024

1 Problem 4

```
[ ]: import sys

import numpy as np
import matplotlib
import matplotlib.pyplot as plt

print(f"Python version: {sys.version}\nNumpy version: {np.
↪ __version__}\nMatplotlib version: {matplotlib.__version__}")
```


Python version: 3.11.8 (main, Feb 26 2024, 15:36:12) [Clang 14.0.6]
Numpy version: 1.26.4
Matplotlib version: 3.8.0

```
[ ]: np.seterr(invalid='ignore', over='ignore') # suppress warning caused by ↪
↪ division by inf

def f(x):
    return 1/(1 + np.exp(3*(x-3))) * 10 * x**2 + 1 / (1 + np.exp(-3*(x-3))) * ↪
↪ (0.5*(x-10)**2 + 50)

def fprime(x):
    return 1 / (1 + np.exp((-3)*(x-3))) * (x-10) + 1/(1 + np.exp(3*(x-3))) * 20 ↪
↪ * x + (3* np.exp(9))/(np.exp(9-1.5*x) + np.exp(1.5*x))**2 * ((0.5*(x-10)**2 ↪
↪ + 50) - 10 * x**2)

x = np.linspace(-5,20,100)
plt.plot(x,f(x), 'k')
plt.show()
```



```
[ ]: def gradient_descent(epochs, alpha):
    theta = np.random.uniform(x.min(), x.max()) # randomly initialises theta
    ↪ in [x.min(), x.max())=[-5,20)

    for _ in range(epochs): # performs GD for a given number of epochs
        theta = theta - alpha * fprime(theta)

    return theta # returns final theta value after GD
```

```
[ ]: def determine_minima(theta, tolerance=1e-2):
    if abs(theta - 0) < tolerance:
        return "sharp"
    elif abs(theta - 10) < tolerance:
        return "wide"
    else:
        raise ValueError(f"theta did not converge to a minimum (within
    ↪ tolerance={1e-2})")
```

```
[ ]: def run_experiment(num_trials, epochs, learning_rate):
    results = {"sharp": 0, "wide": 0, "failed": 0}

    digits = np.floor(np.log10(num_trials)).astype(int) + 1
```

```

for i in range(num_trials):
    print_i = str(i).zfill(digits)
    print(f"Trial num: {print_i}/{num_trials}", end="\r")
    gd_theta = gradient_descent(epochs, learning_rate)
    try:
        results[determine_minima(gd_theta)] += 1
    except ValueError:
        results["failed"] += 1

# results = {k: str(v).zfill(digits) for k, v in results.items()}
results = {k: f"{(v/num_trials) * 100:6.2f}" for k, v in results.items()}
↪# convert results to a percentage

print(f"""In {num_trials} trials (with {epochs} epochs each) using learning_
↪rate alpha = {learning_rate}:
    Convergence to sharp minimum: {results['sharp']}%
    Convergence to wide minimum: {results['wide']}%
Failed to converge to a minimum: {results['failed']}%
""")

```

```

[ ]: num_trials = 150
     epochs = 5000

for alpha in (0.01, 0.30, 4.00):
    run_experiment(num_trials, epochs, alpha)

```

In 150 trials (with 5000 epochs each) using learning rate alpha = 0.01:

Convergence to sharp minimum: 33.33%

Convergence to wide minimum: 66.67%

Failed to converge to a minimum: 0.00%

In 150 trials (with 5000 epochs each) using learning rate alpha = 0.3:

Convergence to sharp minimum: 0.00%

Convergence to wide minimum: 100.00%

Failed to converge to a minimum: 0.00%

In 150 trials (with 5000 epochs each) using learning rate alpha = 4.0:

Convergence to sharp minimum: 0.00%

Convergence to wide minimum: 0.00%

Failed to converge to a minimum: 100.00%

2 Problem 5

```
[ ]: import numpy as np

class Convolution1d :
    def __init__(self, filt) :
        self.__filt = filt
        self.__r = filt.size
        self.T = TransposedConvolution1d(self.__filt)

    def __matmul__(self, vector) :
        r, n = self.__r, vector.size

        return np.asarray([np.dot(
            self.__filt,
            vector[i:i+r]
        ) for i in range(n-r+1)]) # IMPLEMENTATION

class TransposedConvolution1d :
    '''
    Transpose of 1-dimensional convolution operator used for the
    transpose-convolution operation  $A.T@(\dots)$ 
    '''
    def __init__(self, filt) :
        self.__filt = filt
        self.__r = filt.size

    def __matmul__(self, vector) :
        r = self.__r
        n = vector.size + r - 1

        return np.asarray([np.dot(
            np.flip(self.__filt),
            np.concatenate((np.zeros(r-1), vector, np.zeros(r-1)))
        ) for i in range(n)]) # IMPLEMENTATION

def huber_loss(x) :
    return np.sum( (1/2)*(x**2)*(np.abs(x)<=1) + (np.sign(x)*x-1/2)*(np.
↪abs(x)>1) )

def huber_grad(x) :
    return x*(np.abs(x)<=1) + np.sign(x)*(np.abs(x)>1)

r, n, lam = 3, 20, 0.1

np.random.seed(0)
k = np.random.randn(r)
```

```

b = np.random.randn(n-r+1)
A = Convolution1d(k)
#from scipy.linalg import circulant
#A = circulant(np.concatenate((np.flip(k),np.zeros(n-r))))[r-1:,: ]

x = np.zeros(n)
alpha = 0.01
for _ in range(100) :
    x = x - alpha*(A.T@(huber_grad(A@x-b))+lam*x)

print(huber_loss(A@x-b)+0.5*lam*np.linalg.norm(x)**2)

```

0.4587586843129764