

hw12_code_output

June 16, 2024

```
[ ]: import sys
import time

import numpy as np
import matplotlib
import matplotlib.pyplot as plt
import torch

%matplotlib inline
print(f"Python version: {sys.version}\nNumpy version: {np.
↪ __version__}\nMatplotlib version: {matplotlib.__version__}\nPyTorch version:
↪ {torch.__version__}")
```

Python version: 3.11.8 | packaged by conda-forge | (main, Feb 16 2024, 20:49:36)
[Clang 16.0.6]
Numpy version: 1.26.4
Matplotlib version: 3.8.0
PyTorch version: 2.3.1

1 Problem 1

1.1 Setup (given)

```
[ ]: N, p = 30, 20
np.random.seed(0)
X = np.random.randn(N,p)
Y = 2*np.random.randint(2, size = N) - 1
lamda = 30
```

1.2 Training

```
[ ]: theta = 0.1 * np.random.randn(p)
phi = 0.1 * np.random.randn(p)
alpha = 3e-1
beta = 1e-4

epoch = 5000
L_val = []
```

```

d_phi_val = []
d_theta_val = []

for _ in range(epoch):
    for __ in range(N):
        # Use alternating stochastic gradient ascent-descent
        i = np.random.randint(p) # stochastic so choose random index of
        ↪gradient to descend/ascent
        stoc_d_phi = np.average(Y / (1 + np.exp(Y * ((X-phi.reshape(1,-1)) @
        ↪theta)))) * theta[i] - lamda * phi[i]
        phi += beta*stoc_d_phi

        # uses updated phi value
        stoc_d_theta = np.average((-Y / (1 + np.exp(Y * ((X-phi.reshape(1,-1))
        ↪@ theta)))) .reshape(-1,1)*(X-phi[i]), axis=0)
        theta -= alpha*stoc_d_theta

        L_i = np.average(np.log(1 + np.exp(-Y * ((X - phi.reshape(1,-1)) @
        ↪theta)))) - lamda/2 * np.linalg.norm(phi, axis=0, ord=2) **2
        d_phi = np.average(Y / (1 + np.exp(Y * ((X-phi.reshape(1,-1)) @ theta)))) *
        ↪theta - lamda * phi
        d_theta = np.average((-Y / (1 + np.exp(Y * ((X-phi.reshape(1,-1)) @
        ↪theta)))) .reshape(-1,1)*(X-phi.reshape(1,-1)), axis=0)

        L_val.append(L_i)
        d_phi_val.append(d_phi)
        d_theta_val.append(d_theta)

```

1.3 Plot results

```

[ ]: fig, (ax1, ax2, ax3) = plt.subplots(3, 1, figsize=(20, 20))
plt.subplots_adjust(left=0.125,
                    bottom=0.1,
                    right=0.9,
                    top=0.9,
                    wspace=0.2,
                    hspace=0.35)


ax1.set_title("L function")
ax1.plot(L_val)

ax2.set_title(r"$||\nabla_{\theta} L||$")
ax2.plot(np.linalg.norm(d_theta_val, axis=1, ord=2))

ax3.set_title(r"$||\nabla_{\phi} L||$")
ax3.plot(np.linalg.norm(d_phi_val, axis=1, ord=2))

```

```
plt.show()
```



From the graphs, the decreasing $||\nabla_{\theta} L||$ indicates L is being minimised wrt θ while $||\nabla_{\phi} L||$'s higher value and general increase shows L is being maximised wrt ϕ .

The graph of L plateauing after around 1000 epochs suggests that θ and ϕ have reached an equilibrium at this point.

2 Problem 3

2.1 Given setup and dataset

```
[ ]: import torch
import torch.nn as nn
from torch.utils.data import Dataset, TensorDataset, DataLoader
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt

torch.random.manual_seed(0)
batch_size = 64
learning_rate = 5e-4
num_epochs = 2000 # given in exercise
device = "cpu"
# device = "mps" # "cuda:0" if torch.cuda.is_available() else "cpu"

[ ]: def make_swiss_roll(n_samples=2000, noise = 1.0, dimension = 2, a = 20, b = 5):
    """
    Generate 2D swiss roll dataset
    """
    t = 2 * np.pi * np.sqrt(np.random.uniform(0.25,4,n_samples))

    X = 0.1 * t * np.cos(t)
    Y = 0.1 * t * np.sin(t)

    errors = 0.025 * np.random.multivariate_normal(np.zeros(2), np.
eye(dimension), size = n_samples)
    X += errors[:, 0]
    Y += errors[:, 1]
    return np.stack((X, Y)).T

class SwissRollDataset(Dataset) :
    def __init__(self, data) :
        super().__init__()
        self.data = torch.from_numpy(data).to(torch.float32) # to allow for
running on mps

    def __len__(self) :
        return len(self.data)

    def __getitem__(self, idx) :
        return self.data[idx]

def show_data(data, title, levels=1000):
    """
```

```

Plot the data distribution
"""
sns.set(rc={'axes.facecolor': 'honeydew', 'figure.figsize': (5.0, 5.0)})
plt.figure(figsize = (5, 5))
plt.rc('text', usetex = False)
plt.rc('font', family = 'serif')
plt.rc('font', size = 10)

g = sns.kdeplot(x=data[:, 0], y=data[:, 1], fill=True, thresh=0.1,
↪levels=levels, cmap="Greens")


g.grid(False)
plt.margins(0, 0)
plt.xlim(-1.5,1.5)
plt.ylim(-1.5,1.5)
plt.title(title + f' (levels={levels})')
plt.show()

```

```

[ ]: data = make_swiss_roll()
show_data(data, r"$p\_text{data}$")

```



```
[ ]: dataset = SwissRollDataset(data)
loader = DataLoader(dataset, batch_size=batch_size, shuffle=True)
```

2.2 Implement models

```
[ ]: class Encoder(nn.Module):
    def __init__(self, input_dim=2, hidden_width=128):
        super(Encoder, self).__init__()
        self.fc1 = nn.Sequential(nn.Linear(input_dim, hidden_width),
                                   nn.LeakyReLU(0.2))
        self.fc2 = nn.Sequential(nn.Linear(hidden_width, hidden_width),
                                   nn.Tanh())
        self.fc3 = nn.Linear(hidden_width, 2)

    def forward(self, x):
        x = self.fc1(x)
        x = self.fc2(x)
        x = self.fc3(x)
        mu, log_std = x[:,0], x[:,1]
        return mu, log_std

class Decoder(nn.Module):
    def __init__(self, input_dim=1, hidden_width=64):
        super(Decoder, self).__init__()
        self.fc1 = nn.Sequential(nn.Linear(input_dim, hidden_width),
                                   nn.LeakyReLU(0.2))
        self.fc2 = nn.Sequential(nn.Linear(hidden_width, hidden_width),
                                   nn.Tanh())
        self.fc3 = nn.Linear(hidden_width, 2)

    def forward(self, x):
        x = self.fc1(x)
        x = self.fc2(x)
        return self.fc3(x)
```

2.3 Train model

```
[ ]: Encoder = Encoder().to(device)
Decoder = Decoder().to(device)
optimizer = torch.optim.Adam(list(Encoder.parameters()) + list(Decoder.
    ↪parameters()), lr=learning_rate)
F_THETA_SIGMA = 1/torch.sqrt(torch.tensor(150))
```

```
[ ]: # -log_p_theta(x/z) (reconstruction loss)
def log_p(x, mu, log_std, sigma=F_THETA_SIGMA):
    # mu, log_std = Encoder(x)
```

```

    z = mu + torch.exp(log_std * 0.5)*torch.randn_like(mu)  #  $u = \mu + \sigma \epsilon$ 
    ↪ re-parametrisation trick
    mse = nn.MSELoss().to(device)
    return mse(x, Decoder(z.view(z.shape[0], -1)))/torch.square(sigma)

# KL-divergence terms (regularisation loss)
def kl_div(mu, log_std):
    # mu, log_std = Encoder(x)
    return torch.mean(log_std.exp() + mu**2 - log_std)  #  $tr(\sigma^2) + ||\mu||^2 - \log(\det(\sigma^2))$ 
    ↪  $\log(\det(\sigma^2))$ 

```

```

[ ]: def sample_and_show(epoch_num, num_samples=2000):
    z = torch.randn((num_samples,)).to(device)
    f_theta = Decoder(z.view(z.shape[0], -1))
    x = torch.normal(f_theta, F_THETA_SIGMA * torch.ones_like(f_theta)).cpu()

    show_data(x.detach().numpy(), f"Epoch: {epoch_num}")

for epoch in range(num_epochs):
    # Visualize the intermediate result
    if epoch % (num_epochs // 5) == 0:
        sample_and_show(epoch)

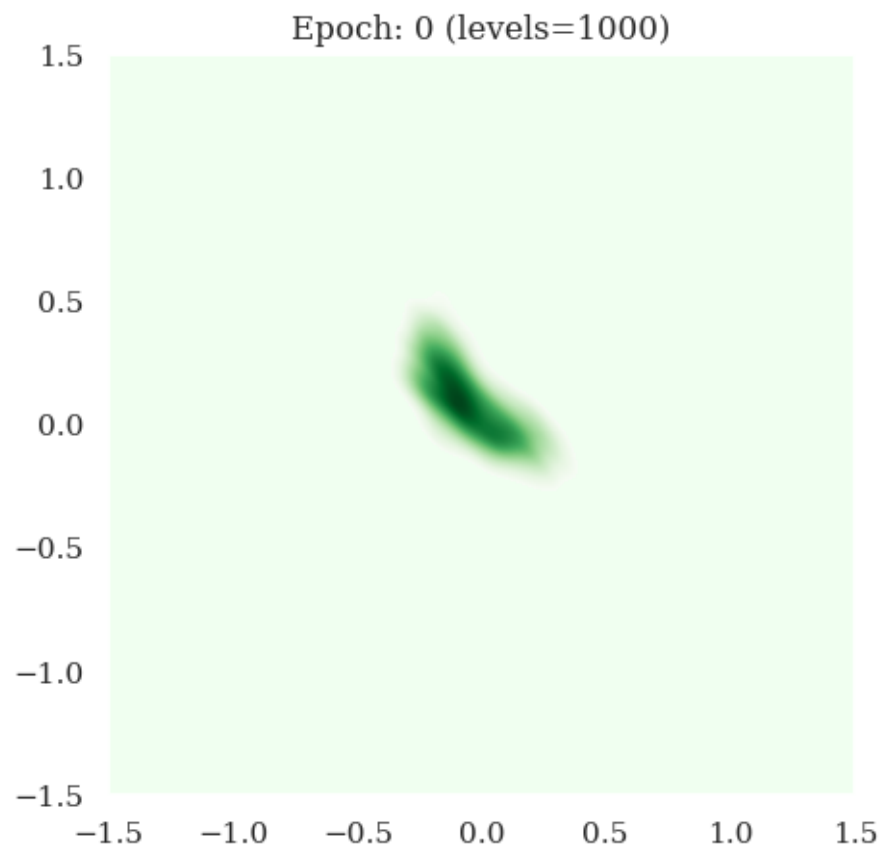
    for x in loader:
        x = x.to(device)
        mu, log_std = Encoder(x)

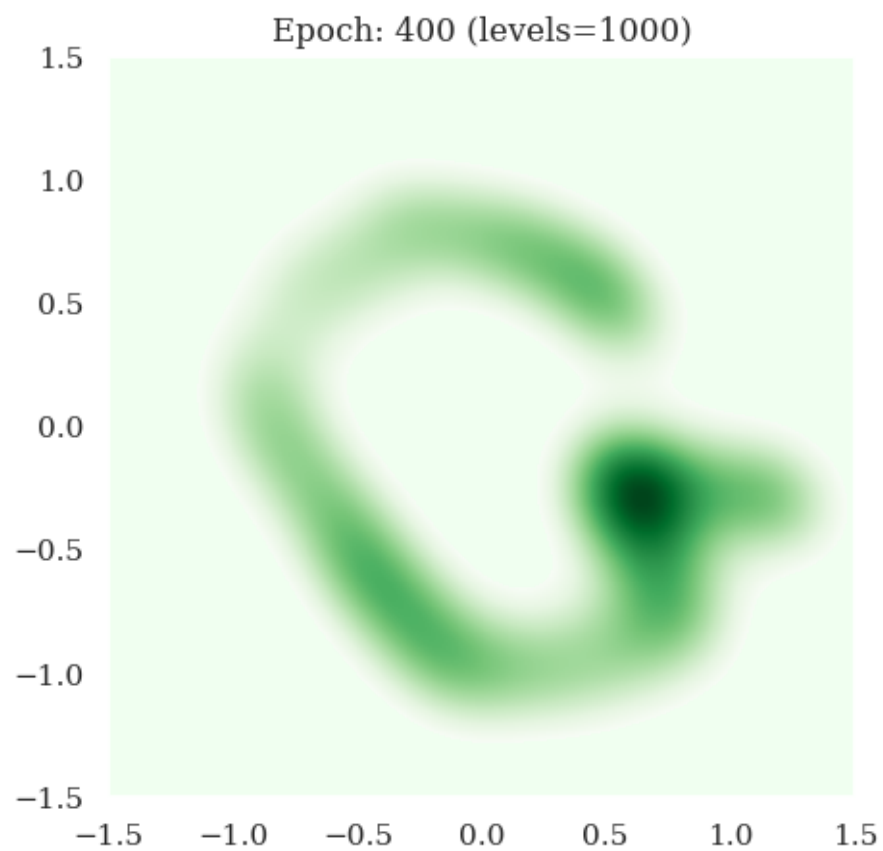
        loss = log_p(x, mu, log_std) + kl_div(mu, log_std)


        loss.backward()
        optimizer.step()
        optimizer.zero_grad()

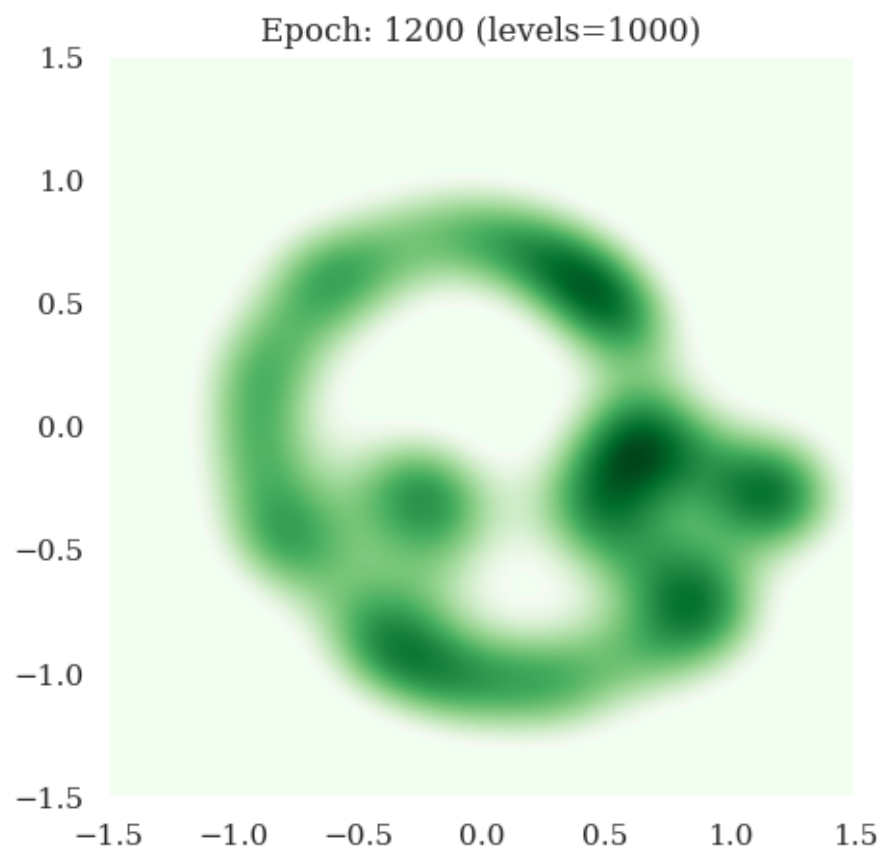
sample_and_show(num_epochs)
show_data(data, r"True Data ($p\_text{data}$)")


```




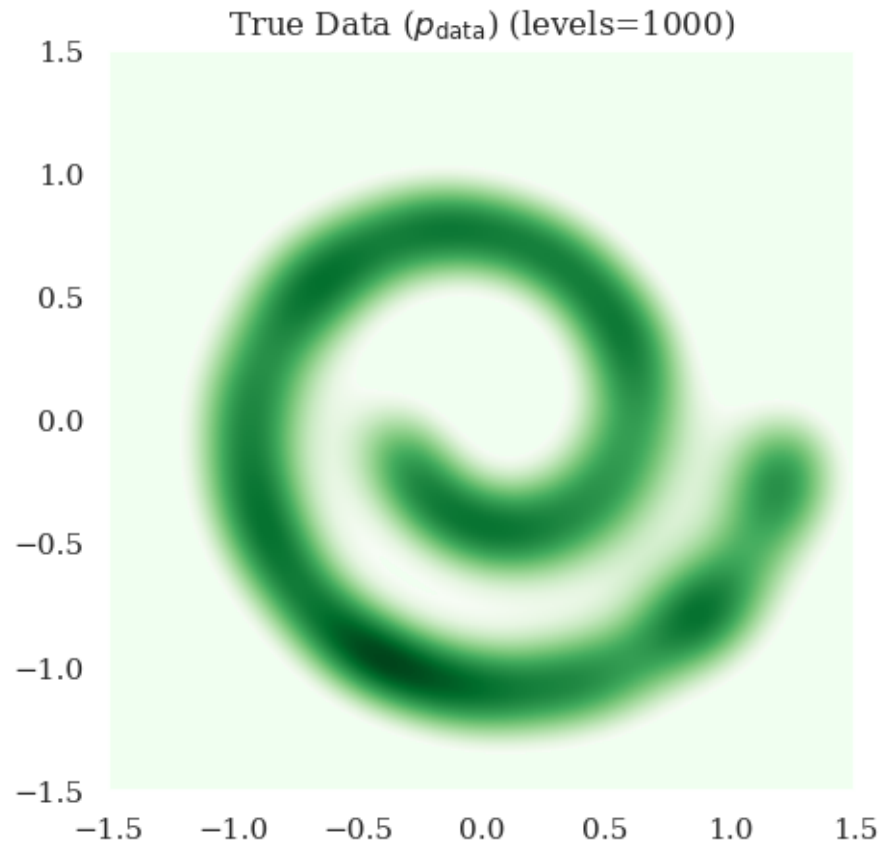












3 Problem 4

```
[ ]: torch.random.manual_seed(0)
num_epochs = 2000
dataset = SwissRollDataset(make_swiss_roll())
loader = DataLoader(dataset, batch_size=batch_size, shuffle=True)
```

3.1 Implement models

```
[ ]: class Discriminator(nn.Module):
    def __init__(self, input_dim=2, hidden_width=128):
        super(Discriminator, self).__init__()
        self.fc1 = nn.Sequential(nn.Linear(input_dim, hidden_width),
                                   nn.Tanh())
        self.fc2 = nn.Sequential(nn.Linear(hidden_width, hidden_width),
                                   nn.Tanh())
        self.fc3 = nn.Sequential(nn.Linear(hidden_width, 1),
                                   nn.Sigmoid())
```

```

def forward(self, x):
    x = self.fc1(x)
    x = self.fc2(x)
    x = self.fc3(x)
    return x

class Generator(nn.Module):
    def __init__(self, input_dim=1, hidden_width=32):
        super(Generator, self).__init__()
        self.fc1 = nn.Sequential(nn.Linear(input_dim, hidden_width),
                                   nn.Tanh())
        self.fc2 = nn.Linear(hidden_width, 2)

    def forward(self, x):
        x = self.fc1(x)
        x = self.fc2(x)
        return x

```

3.2 Train model

```

[ ]: D = Discriminator().to(device)
     G = Generator().to(device)

D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate)
G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate)

```

```

[ ]: def sample_and_show(epoch_num, num_samples=2000):
        z = torch.randn((num_samples, 1)).to(device)
        generated = G(z)
        levels = 1000
        while levels > 0:
            try:
                # data starts off much more bunched so fewer levels required to
                avoid "Contour levels must be increasing" error
                show_data(generated.detach().numpy(), f"Epoch: {epoch_num}",
                levels=levels)
            except ValueError:
                levels -= 100
            finally:
                return

        raise ValueError("Failed to find an appropriate levels value for sns.
        ↳ kdeplot that allowed data to be drawn")

for epoch in range(num_epochs):
    # Visualize the intermediate result
    if epoch % (num_epochs // 5) == 0:

```

```

sample_and_show(epoch)

for x in loader:
    x.to(device)
    # use alternating ascent descent GD
    D.zero_grad()
    D_type_1_loss = torch.mean(torch.log(D(x)))
    z = torch.randn((batch_size, 1)).to(device)
    D_type_2_loss = torch.mean(torch.log(1-D(G(z))))

    D_loss = -(D_type_1_loss + D_type_2_loss) # maximise
    D_loss.backward()
    D_optimizer.step()

    G.zero_grad()
    z = torch.randn((batch_size, 1)).to(device)
    G_loss = torch.mean(torch.log(1-D(G(z)))) # minimise
    G_loss.backward()
    G_optimizer.step()

sample_and_show(num_epochs)
show_data(data, r"True Data ($p_{\text{data}}$)")

```